

## Model-free reinforcement learning: Q-learning and SARSA

Model-free reinforcement learning: what if we do not know transitions  $P$  and reward function  $r$  of an MDP?

The Mystery Game:

<https://programmingheroes.blogspot.com/2016/02/udacity-reinforcement-learning-mystery-game.html>

### Model-based vs model-free

MDP:

- Set of states  $S$
- Transition probabilities  $P_{a(s' / s)}$
- Reward function  $r(s, a, s')$  in real
- Discount factor  $\gamma$

### Q-learning

1. Initialise  $Q(s,a)$  arbitrarily
2. For each episode:
  - a. Initialise  $s$  (go to the initial state)
  - b. Repeat for each step in the episode
    - i. Select the next action  $a$  to apply from  $s$  (using e.g. epsilon greedy, UCT) using  $Q(s,a)$
    - ii. Execute action  $a$  and observe the reward  $r$  and new state  $s'$
    - iii.  $Q(s,a) := Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$
    - iv.  $s := s'$
  - c. Until  $s$  is terminal

$$Q(s,a) := Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$


Handwritten annotations:

- $(1-\alpha)$  reward: points to the  $Q(s,a)$  term on the left.
- estimated reward: points to the  $r$  term.
- estimated future reward: points to the  $\gamma \max_{a'} Q(s',a')$  term.
- learning rate  $0 < \alpha < 1$ : points to the  $\alpha$  term.
- discount factor: points to the  $\gamma$  term.

### Q-Tables

State	Action			
	North	South	East	West
(0,0)	0.53	0.36	0.36	0.21
(0,1)	0.61	0.27	0.23	0.23
...				
(2,2)	0.79	0.72	0.90	0.72
(2,3)	0.90	0.78	0.99	0.81

## SARSA example



State	Action			
	North	South	East	West
(0,0)	0.53	0.36	0.36	0.21
(0,1)	0.61	0.27	0.23	0.23
...				
(2,2)	0.79	0.72	0.90	0.72
(2,3)	0.90	0.78	0.99	0.81

$$Q(s,a) := Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$

Learning rate  $\alpha = 0.1$

Discount reward factor  $\gamma = 0.9$

$$\text{Q-learning: } Q((2,2), \text{North}) = 0.79 + 0.1 * (0 + 0.9 * 0.9 - 0.79) = 0.792$$

$$Q(s,a) := Q(s,a) + \alpha[r + \gamma Q(s',a') - Q(s,a)]$$

SARSA, with assumption that  $a'$  is West (on-policy action)

$$Q((2,2), \text{North}) = 0.79 + 0.1 * (0 + 0.9 * 0.72 - 0.79) = 0.7758$$

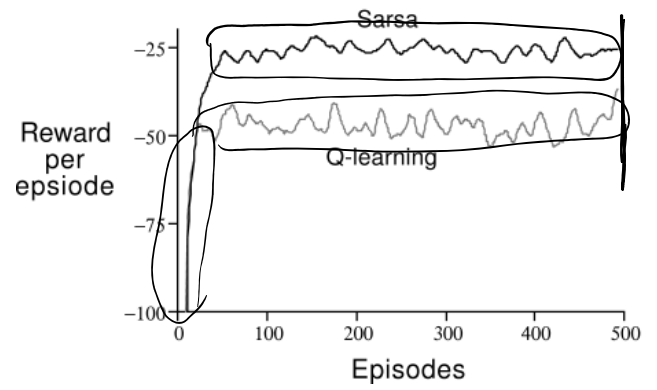

If  $a'$  was East is for SARSA, the update would be just the same as for Q-learning because East is the max action from (2,2)

## Q-learning: Off-policy

1. Initialise  $Q(s,a)$  arbitrarily
2. For each episode:
  - a. Initialise  $s$  (go to the initial state)
  - b. Repeat for each step in the episode
    - i. Select the next action  $a$  to apply from  $s$  (using e.g. epsilon greedy, UCT) use  $Q(s,a)$
    - ii. Execute action  $a$  and observe the reward  $r$  and new state  $s'$
    - iii.  $Q(s,a) := Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$
    - iv.  $s := s'$
  - c. Until  $s$  is terminal

## SARSA: On-policy learning

1. Initialise  $Q(s,a)$  arbitrarily
2. For each episode:
  - a. Initialise  $s$  (go to the initial state)
  - b. Select the next action  $a$  to apply from  $s$  (using e.g. epsilon greedy, UCT)
  - c. Repeat for each step in the episode
    - i. Execute action  $a$  and observe the reward  $r$  and new state  $s'$
    - ii. Select the next action  $a'$  to apply from  $s'$  (using e.g. epsilon greedy, UCT)
    - iii.  $Q(s,a) := Q(s,a) + \alpha[r + \gamma Q(s',a') - Q(s,a)]$
    - iv.  $s := s'; a := a'$
  - d. Until  $s$  is terminal



[Gridworld Q-Learning - Example 3 - The Cliff](#)



[Learning to Play Freeway, using Reinforcement Learning](#)